

Practical No. 11

Aim: To perform and Data analysis with Confusion matrix

In [3]:

```
#Name : Taufiq Rafik Nagori  
#Roll no. : 77 (BDA-B77)  
#Section : B  
#Subject : PE-II
```

In [5]:

```
import os  
import pandas as pd  
import numpy as np
```

In [7]:

```
os.getcwd()
```

Out[7]:

```
'C:\\Users\\USER'
```

In [9]:

```
os.chdir("C:\\Users\\USER\\Desktop")
```

In [11]:

```
data=pd.read_csv("heart.csv")
```

In [13]:

```
data.head()
```

Out[13]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0

In [15]:

```
data.tail()
```

Out[15]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
1020	59	1	1	140	221	0	1	164	1	0.0	2	0	2	1
1021	60	1	0	125	258	0	0	141	1	2.8	1	1	3	0
1022	47	1	0	110	275	0	0	118	1	1.0	1	1	2	0
1023	50	0	0	110	254	0	0	159	0	0.0	2	0	2	1
1024	54	1	0	120	188	0	1	113	0	1.4	1	1	3	0

In [17]:

```
data.info
```

Out[17]:

```
<bound method DataFrame.info of
h  exang  oldpeak  \
0      52      1      0      125      212      0      1      168      0      1.0
1      53      1      0      140      203      1      0      155      1      3.1
2      70      1      0      145      174      0      1      125      1      2.6
3      61      1      0      148      203      0      1      161      0      0.0
4      62      0      0      138      294      1      1      106      0      1.9
...    ...    ...    ..    ...    ...    ...    ...    ...    ...    ...
1020    59      1      1      140      221      0      1      164      1      0.0
1021    60      1      0      125      258      0      0      141      1      2.8
1022    47      1      0      110      275      0      0      118      1      1.0
1023    50      0      0      110      254      0      0      159      0      0.0
1024    54      1      0      120      188      0      1      113      0      1.4

      slope  ca  thal  target
0          2  2    3        0
1          0  0    3        0
2          0  0    3        0
3          2  1    3        0
4          1  3    2        0
...    ...  ..    ...    ...
1020      2  0    2        1
1021      1  1    3        0
1022      1  1    2        0
1023      2  0    2        1
1024      1  1    3        0
```

[1025 rows x 14 columns]>

In [19]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   age         1025 non-null   int64
1   sex         1025 non-null   int64
2   cp          1025 non-null   int64
3   trestbps    1025 non-null   int64
4   chol        1025 non-null   int64
5   fbs         1025 non-null   int64
6   restecg     1025 non-null   int64
7   thalach     1025 non-null   int64
8   exang       1025 non-null   int64
9   oldpeak     1025 non-null   float64
10  slope       1025 non-null   int64
11  ca          1025 non-null   int64
12  thal        1025 non-null   int64
13  target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

In [21]:

```
data.describe()
```

Out[21]:

	age	sex	cp	trestbps	chol	fbs	restecg	
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	102
mean	54.434146	0.695610	0.942439	131.611707	246.000000	0.149268	0.529756	14
std	9.072290	0.460373	1.029641	17.516718	51.59251	0.356527	0.527878	2
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	7
25%	48.000000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	13
50%	56.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	15
75%	61.000000	1.000000	2.000000	140.000000	275.000000	0.000000	1.000000	16
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	20

In [23]:

```
data.shape
```

Out[23]:

(1025, 14)

In [25]:

```
data.size
```

Out[25]:

14350

In [27]:

```
data.ndim
```

Out[27]:

2

In [29]:

```
data.columns
```

Out[29]:

```
Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',  
      'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],  
      dtype='object')
```

Data pre-processing, data-cleaning, missing value treatment

In [32]:

```
data.isna()
```

Out[32]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	False	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False	False

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
4	False	False	False	False	False	False	False	False	False	False	False	False	False
...
1020	False	False	False	False	False	False	False	False	False	False	False	False	False
1021	False	False	False	False	False	False	False	False	False	False	False	False	False
1022	False	False	False	False	False	False	False	False	False	False	False	False	False
1023	False	False	False	False	False	False	False	False	False	False	False	False	False
1024	False	False	False	False	False	False	False	False	False	False	False	False	False

1025 rows × 14 columns

In [34]:

```
data.isna().any()
```

Out[34]:

```
age      False
sex      False
cp       False
trestbps False
chol     False
fbs      False
restecg  False
thalach  False
exang    False
oldpeak  False
slope    False
ca       False
thal     False
target   False
dtype: bool
```

In [36]:

```
data.isna().sum()
```

Out[36]:

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

Remove Duplicates

In [39]:

```
data_dup =data.duplicated().any()
```

In [41]:

```
data_dup
```

Out[41]:

True

In [43]:

```
data=data.drop_duplicates()
```

In [45]:

```
data_dup =data.duplicated().any()
```

In [47]:

```
data_dup
```

Out[47]:

False

Splitting dataset into training and testing

In [50]:

```
x = data.drop('target', axis=1)
y = data['target']
```

In [52]:

```
from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2 ,random_state=42)
```

In [54]:

```
x_train
```

Out[54]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
163	48	1	0	124	274	0	0	166	0	0.5	1	0	3
291	58	1	0	128	259	0	0	130	1	3.0	1	2	3
280	45	0	1	130	234	0	0	175	0	0.6	1	0	2
85	44	1	1	120	220	0	1	170	0	0.0	2	0	2
239	62	0	0	150	244	0	1	154	1	1.4	1	0	2
...
267	67	1	0	120	237	0	1	71	0	1.0	1	0	2
77	63	1	0	140	187	0	0	144	1	4.0	2	2	3
125	60	0	3	150	240	0	1	171	0	0.9	2	0	2
522	67	0	2	152	277	0	1	172	0	0.0	2	1	2
119	42	1	1	120	295	0	1	162	0	0.0	2	0	2

241 rows × 13 columns

In [56]:

```
x_test
```

Out[56]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
245	44	1	1	130	219	0	0	188	0	0.0	2	0	2
349	62	0	2	130	263	0	1	97	0	1.2	1	1	3
135	58	0	0	170	225	1	0	146	1	2.8	1	2	1
389	63	1	3	145	233	1	0	150	0	2.3	0	0	1
66	53	1	2	130	197	1	0	152	0	1.2	0	0	2
...
402	70	1	1	156	245	0	0	143	0	0.0	2	0	2
123	65	0	2	140	417	1	0	157	0	0.8	2	1	2
739	52	1	0	128	255	0	1	161	1	0.0	2	1	3
274	66	1	0	160	228	0	0	138	0	2.3	2	0	1
256	35	0	0	138	183	0	1	182	0	1.4	2	0	2

61 rows × 13 columns

In [58]:

```
y_train
```

Out[58]:

```
163    0
291    0
280    1
85     1
239    0
..
267    0
77     0
125    1
522    1
119    1
```

Name: target, Length: 241, dtype: int64

In [60]:

```
y_test
```

Out[60]:

```
245    1
349    0
135    0
389    1
66     1
..
402    1
123    1
739    0
274    1
```

256 1
Name: target, Length: 61, dtype: int64

Logistic Regression

In [63]:

```
from sklearn.linear_model import LogisticRegression
```

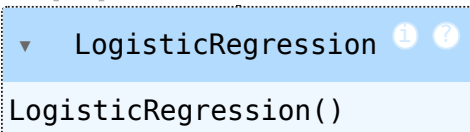
In [65]:

```
log = LogisticRegression()  
log.fit(x_train, y_train)
```

C:\Users\USER\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

Out[65]:



LogisticRegression()

In [67]:

```
y_predict=log.predict(x_test)
```

In [69]:

```
from sklearn.metrics import accuracy_score  
accuracy_score (y_test,y_predict)
```

Out[69]:
0.8032786885245902

Confusion Matrix

In [72]:

```
from sklearn.naive_bayes import GaussianNB
```

In [74]:

```
from sklearn.metrics import confusion_matrix, classification_report
```

In [76]:

```
model = GaussianNB()  
model.fit(x_train, y_train)  
y_predict = model.predict(x_test)
```

In [78]:


```
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.metrics import confusion_matrix
```

In [80]:

```
confusionMatrix = confusion_matrix(y_test, y_predict)
```

```
plt.figure(figsize=(6,4))

sns.heatmap(confusionMatrix, annot=True, fmt = 'd', cmap="Reds", linewidths=1, linecolor='b'
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



In [82]:

```
print("Classification Report: \n")
print(classification_report(y_test,y_predict))
```

Classification Report:

	precision	recall	f1-score	support
0	0.87	0.84	0.86	32
1	0.83	0.86	0.85	29
accuracy			0.85	61
macro avg	0.85	0.85	0.85	61
weighted avg	0.85	0.85	0.85	61